BuildMap — work completed, remaining work and additions

Status at the saved checkpoint • October 21, 2026

BuildMap implementation was delivered and pushed. The whole-code scan completed. The sequential end-to-end workflow review remains unfinished.

Your original request

Review the current BuildMap repository against BuildMap_Codex_Project_Roadmap_20261021.docx, classify its features, and prioritize Calculation & Data Flow Inspector and Pre-Run Risk Review. Use Frankie as the current test case, keep BuildMap general-purpose, and defer outside-customer work.

What I completed

• Reviewed the actual application baseline and recorded Existing / Partial / Missing statuses for all 12 roadmap items before implementation.

• Added Python calculation inspection, source navigation, candidate cross-file calls, possible producer/consumer products and evidence-linked pre-run risks.

• Added checkpoint/resource review, contract-based row/column extension plans, declared-stage checks and optional reported-workload comparisons.

• Preserved the existing free map, evidence states, walkthroughs and task briefs. Replaced superseded check records; no evidence-history subsystem was added.

• Committed and pushed the implementation in the BuildMap repository on the separate workflow-inspector branch after you confirmed the other audit was done.

What I added after your follow-up requests

• Scan-coverage panel: shows unreviewed Python files and inventory files outside the Python risk adapter.

• Runtime-to-calculation links: reported records can open a captured calculation; stale or ambiguous references cannot select a target or justify a comparison.

• Whole-project snapshots: --workflow-all processes every eligible Python file in bounded batches. The 64-file total cutoff is removed in this mode; the default live preview remains bounded.

Verification: 143/143 Node tests and 14/14 scoped Chromium checks passed. These checks do not establish a completed Frankie production workflow or renew the older whole-MVP completion claim.

Roadmap accounting — BM-01 to BM-06

Baseline means the application before this work. Current status is Partial for every full roadmap item: useful implementation exists, but the complete acceptance scope is not finished. No broad item is claimed fully Existing.

Roadmap item	Before to now	Implemented / preserved	Still not completed
BM-01 Calculation & Data Flow Inspector	Partial to Partial	Python modules/functions/direct methods; inputs, assignments, returns, source hashes/locations and candidate source calls.	Full semantic data propagation, dynamic dispatch, business-purpose/formula verification and complete language coverage.
BM-02 Pre-Run Architecture / Workload Risk Review	Missing to Partial	Review cues for repeated reads, loop expansion, memory retention, possible waits and saved-state compatibility; actions and verification guidance.	A proven defect list, measured impact and automatic clearance for the entire workflow.
BM-03 Producer–Consumer & Reuse Map	Partial to Partial	Possible shared literal-path products; producers/consumers and supplied exact identity/coverage contracts.	Complete real product lineage, equivalent-producer detection and verified shared reuse across dynamic paths.
BM-04 Workload / Combinatorial Explosion	Missing to Partial	Symbolic loop/comprehension expansion; exact literal-range potential counts; jobs and directional pairs kept separate.	Actual remaining jobs/pairs, supported runtime forecasts and complete SQL/Cartesian expansion analysis.
BM-05 Duplicate Computation / Data Reads	Missing to Partial	Repeated literal read sites, whole-result construction and existing cache/lazy/FFT/index/restart cues for review.	Proof of duplicated execution, equivalent computation detection and measured avoidable work/savings.
BM-06 Checkpoint & Restart Paths	Missing to Partial	Persistence/helper sites, branch guards and nearby existence/identity checks; incompatible reuse remains unverified.	Complete durable-owner/recovery semantics, demonstrated restart equivalence and certified safe resume.

These statuses distinguish delivered implementation from full-roadmap completion. Static findings, conditional contracts and source-call candidates do not prove execution, scientific correctness or safe reuse.

Roadmap accounting — BM-07 to BM-12

Baseline means the application before this work. Current status is Partial for every full roadmap item: useful implementation exists, but the complete acceptance scope is not finished. No broad item is claimed fully Existing.

Roadmap item	Before to now	Implemented / preserved	Still not completed
BM-07 Incremental Table Coverage / Extension	Missing to Partial	Exact per-column interval unions, holes, reusable/missing cells and conditional extension plans; declared rolling overlap; stateful extension blocked.	Actual scientific-product compatibility checks, domain-specific state boundaries and executing a correct extension.
BM-08 CPU / Memory / I/O / Parallelism	Missing to Partial	Allocation/retention and wait cues; pool/chunk/release controls; optional supplied CPU/memory/I/O/time metrics.	Live profiling, full inter-host topology and a measured explanation of each production bottleneck.
BM-09 Intended vs Actual / Agent Drift	Partial to Partial	Existing bounded tasks retained; declared stage references, cycles, competing producers and ordering questions checked.	Complete semantic comparison between user intent and every executable workflow path.
BM-10 Guided Walkthrough + Task Briefs	Partial to Partial	Existing free walkthroughs/briefs retained; calculation reading and source navigation extended within the Guide preview.	A complete novice-friendly project tour, durable coaching continuity and commercial entitlement/billing.
BM-11 Runtime Evidence / Savings	Missing to Partial	Optional supplied counters/metrics and receipt freshness; comparisons require matching identities/scope; exact signed differences.	Automatic live telemetry/log integration, genuine Frankie runtime overlay and validated optimization savings.
BM-12 Frankie Dogfood / Regression Pack	Missing to Partial	Small generic regression fixtures plus a real current AWS whole-code scan; existing optimizations inspected; source hash checks retained.	Completed sequential end-to-end handoff review and demonstrated full scientific/production workflow correctness.

These statuses distinguish delivered implementation from full-roadmap completion. Static findings, conditional contracts and source-call candidates do not prove execution, scientific correctness or safe reuse.

Whole-code run and unfinished workflow review

The current AWS checkout was accessed after you supplied the recent branch. The final scan did not use the October 6 or October 14 checkout as current evidence.

Observed scope / result	Recorded value
Current source branch	codex/search-assembly-20261021
Stable source commit during the full scan	ad00e5b54814ac12bf6a4f8662182d6b02869834
Entire inventoried code/workflow set	1,856 files
Python source parsed / skipped	1,469 parsed; 0 Python files skipped
Other inventoried files	387; execution semantics are outside the Python risk adapter
Calculation records / candidate source calls	18,577 / 29,363
Static review cues / unsupported call links	18,857 cues; 192,236 call sites lack a supported source link
Execution environment	Node 20.20.2 on AWS; target workflow was not launched

The 18,857 cues are review questions, not 18,857 confirmed bugs. Parsing every Python file is not proof that every dynamic path or end-to-end handoff works.

What I did not finish

• The final, prioritized review of the complete workflow from intake through reports, successor handling and owner/day closure.

• Runtime validation that every producer hands the correct complete product to every consumer, including all guarded newer indexed paths.

• Fresh, valid native owner/assembly liveness verification. One later status read exceeded the SSM output limit; its truncated result was not accepted as evidence.

• Measured Frankie original/eligible/completed/remaining jobs, directional pairs, actual reads, cache effectiveness, ETA or savings.

• Scientific parity, full restart/continuation correctness, all-language execution analysis or a completed production end-to-end run.

Some queue paths and guarded consumer boundaries were already inspected, but the complete sequential review and final verdict were not delivered. Old saved failure receipts were not treated as proof of current failure.

Scope boundaries, additions and continuation

Actions and boundaries

• All committed BuildMap product changes are in DavisAI1974/BuildMap. Temporary extra local BuildMap worktrees were consolidated into /workspace/BuildMap. No BuildMap implementation branch was created in Markets.

• The initial obsolete-source investigation was superseded. The full-code scan used the user-identified current AWS branch and a stable newer commit; its source and deployed/saved runtime identities were kept separate.

• To run the full scan, I created a separate BuildMap Git checkout on AWS and installed its pinned JavaScript dependency there. The scan and private outputs stayed in the BuildMap repository, not in the Frankie checkout.

• I did not edit Frankie source, mutate its data/indexes/checkpoints, launch/resume/stop its workflow, start a 30-day run, re-ingest the large dataset or change its infrastructure.

• I did not merge the BuildMap branch, deploy a customer-facing product, add accounts/billing or build an evidence-history/lineage subsystem. Optional before/after runtime comparisons use only records the user supplies.

Continue the workflow review in these pieces

1. Intake/fetch to sealed ingest to external/day attachment to ROOT admission.

2. ROOT outputs to teacher/hub prerequisites to classroom ownership.

3. Shared preparation to classroom and independent Search/calculation branches to their joins and restart/cache reuse.

4. Batch scientific lessons to accumulated/prior claims to novel Frankie findings to survivor updates.

5. Exchange to Jev and checked knowledge to voice/meeting to school to reports.

6. Late report revisions to successor drain to final owner/day closure; check save/resume/failure handoffs throughout.

Delivery and restart pointers

BuildMap branch: private/workflow-inspector-risk-20261021. Implementation checkpoint: e6a30c7. Saved review handoff: a374113. These checkpoints were pushed; the implementation branch was not merged.

Read docs/REVIEW_RESUME_20261021.md first in the next chat. Private full scan artifacts remain in /opt/frankie-box/work/BuildMap/.git/whole-workflow-review/ on AWS. Local private inspection material is under /workspace/BuildMap/.git/buildmap-review/ if the workspace persists.

Restart prompt: “Continue BuildMap’s complete workflow review on branch private/workflow-inspector-risk-20261021. Read docs/REVIEW_RESUME_20261021.md. Continue the sequential end-to-end review from the checkpoint; do not repeat the completed scan unless relevant source changes justify it.”